

- Maintaining independence from commercial AI providers and their pricing models or usage restrictions.

Open source models like BLOOM, GPT-Neo, and Llama provide flexible alternatives to proprietary models while requiring technical expertise for setup and maintenance. When deciding

between open source and proprietary solutions, key differences must be considered (Table 1).

WHAT YOU SHOULD CONSIDER BEFORE BUILDING YOUR OWN LLM

Here's a detailed breakdown of crucial factors.

Define your use case

Specificity. Don't just say 'chatbot'. Define the exact purpose. Will it be for customer service, technical support, creative writing assistance, code generation, or something else? The more specific you are, the better you can tailor your LLM.

- **Target audience.** Who will be using this LLM? Their needs and expectations will influence the model's design and training.
- **Desired output.** What kind of output are you looking for? Text, code, summaries, translations? The format and style of the output are important considerations.
- **Performance metrics.** How will you measure the success of your LLM? Accuracy, fluency, relevance, response time? Define these metrics before you start building.
- **Integration.** How will the LLM integrate with your existing systems and workflows? APIs, deployment environment, etc.
- **Ethical considerations.** What are the potential biases in your data and how will you mitigate them? How will you ensure responsible use of the LLM?

Hardware requirements

Training an LLM is computation-intensive and demands robust

hardware. Consider:

GPUs/TPUs. Nvidia A100, H100, or Google TPUs are commonly used for deep learning tasks.

Memory requirements. Higher RAM and VRAM are needed to handle large datasets efficiently.

Cloud vs on-premises. AWS, Google Cloud, and Azure offer scalable GPU/TPU instances. More control over infrastructure requires a higher upfront investment.

Data: The training material

Data is very crucial to any LLM. It's what fuels its learning. First, you need to find your data: public datasets, web scraping, or your own files. But quantity isn't everything—quality is key. Clean, consistent, and unbiased data is essential. Think of it as giving your LLM a healthy diet. A small amount of good data is better than mountains of junk. Make sure your data is diverse, too, reflecting the real world.

Before feeding it to your LLM, you'll need to prepare it: breaking down text, ensuring consistency, and handling missing bits. And if you're using sensitive data, protect it! Privacy is paramount.

Budget and regulatory compliance

Building an LLM involves two crucial, intertwined aspects: your budget and regulatory compliance. Think of them as the financial and legal guardrails of your project.

First, you need a realistic budget. This covers everything from cloud computing costs (which can be substantial) and data acquisition fees to salaries for your team of experts. Don't forget software, tools, and ongoing maintenance. But your budget isn't just about spending; it's also about spending wisely. That's where regulatory compliance




Fig. 1: Text processing pipeline

TABLE 1

KEY DIFFERENCES BETWEEN OPEN SOURCE AND PROPRIETARY SOLUTIONS

Feature	Open Source LLMs	Proprietary LLMs
Customisation	Full control over modifications and fine-tuning	Limited customisation options
Cost	Typically lower upfront cost but require infrastructure	Pay-per-use pricing is potentially expensive
Data privacy	Can be deployed on-premises for complete control	Data may be stored and processed externally
Scalability	Require own infrastructure and expertise	Scalable cloud-based solutions
Support and updates	Community-driven support	Enterprise-grade support and frequent updates